

000 001 SCRAPL: SCATTERING TRANSFORM WITH RANDOM 002 PATHS FOR MACHINE LEARNING 003 004

005 **Anonymous authors**
006 Paper under double-blind review

007 008 ABSTRACT 009

011 The Euclidean distance between wavelet scattering transform coefficients (known
012 as *paths*) provides informative gradients for perceptual quality assessment of deep
013 inverse problems in computer vision, speech, and audio processing. However, these
014 transforms are computationally expensive when employed as differentiable loss
015 functions for stochastic gradient descent due to their numerous paths, which signifi-
016 cantly limits their use in neural network training. Against this problem, we propose
017 “Scattering transform with Random Paths for machine Learning” (SCRAPL): a
018 stochastic optimization scheme for efficient evaluation of multivariable scattering
019 transforms. We implement SCRAPL for the joint time–frequency scattering trans-
020 form (JTFS) which demodulates spectrotemporal patterns at multiple scales and
021 rates, allowing a fine characterization of intermittent auditory textures. We apply
022 SCRAPL to differentiable digital signal processing (DDSP), specifically, unsu-
023 pervised sound matching of a granular synthesizer and the Roland TR-808 drum
024 machine. We also propose an initialization heuristic based on importance sampling,
025 which adapts SCRAPL to the perceptual content of the dataset, improving neural
026 network convergence and evaluation performance. We make our audio samples
027 available and provide SCRAPL as a Python package.

028 1 INTRODUCTION 029

030 A scattering transform (ST) is a wavelet-based nonlinear operator which decomposes a high-resolution
031 input \mathbf{x} into a collection $\Phi\mathbf{x}$ of low-resolution coefficients, known as *paths* (Mallat, 2012). Without
032 loss of generality, let us consider a two-layer multivariable ST of a time-domain signal $\mathbf{x}(t)$:

$$034 \quad \Phi\mathbf{x}(p, t, \lambda) = \rho \left(\left(|\mathbf{W}\mathbf{x}| \circledast \Psi_p \right) \circledast \Psi_0 \right)(t, \lambda). \quad (1)$$

036 In the equation above, \mathbf{W} is a wavelet transform; the vertical bars denote complex modulus; the circled
037 asterisk \circledast denotes a multivariable convolution over time t and wavelet scale λ ; Ψ is a multivariable
038 wavelet filterbank which is indexed by path p ; Ψ_0 , i.e., Ψ_p with $p = 0$ is a multivariable low-pass
039 filter; and ρ is a pointwise nonlinearity, e.g., path normalization and logarithmic transformation.

040 The design of the filterbank Ψ aims at a tradeoff between three properties: invariance to rigid
041 motion, stability to small deformations, and separation of sparse patterns (Mallat, 2016). In speech
042 and audio processing, examples of such Ψ include “plain” time ST (Andén & Mallat, 2014); joint
043 time–frequency scattering (JTFS) (Andén & Mallat, 2014); and spiral ST (Lostanlen & Mallat, 2016).
044 In computer vision, examples include “plain” 2-D ST (Bruna & Mallat, 2013); joint roto-translation
045 ST Sifre & Mallat (2013); and scalo-roto-translation ST (Oyallon et al., 2014).

046 The squared Euclidean distance between scattering coefficients, or *ST distance* for short, is:
047

$$048 \quad d_\Phi(\mathbf{x}, \tilde{\mathbf{x}}) = \sum_{p=0}^{P-1} \sum_{t=0}^{T-1} \sum_{\lambda=0}^{\Lambda-1} \left| \Phi\mathbf{x}(p, t, \lambda) - \Phi\tilde{\mathbf{x}}(p, t, \lambda) \right|^2, \quad (2)$$

051 where P is the number of paths; T is the number of time samples; and Λ is the number of scales.
052 Behavioral studies suggest that ST distance is a good predictor of dissimilarity judgments between
053 isolated sounds, for suitably chosen Ψ and ρ (Patil et al., 2012; Lostanlen et al., 2021; Tian et al.,
2025). Relatedly, neurophysiology studies suggest that JTFS is a suitable idealized model of

054 spectrotemporal receptive fields in the auditory cortex of humans (Norman-Haignere & McDermott,
 055 2018) and nonhuman mammals (Kowalski et al., 1996). These findings motivate the use of JTFS as
 056 part of a differentiable loss function for neural audio models (Vahidi et al., 2023).

057 As an illustration, let \mathbf{x} be a fixed reference and $\tilde{\mathbf{x}} = F_{\mathbf{x}}(\mathbf{w})$ be its reconstruction by an autoencoder
 058 F with trainable weights \mathbf{w} . Denoting the set of path indices by $\mathcal{P} = \{0, \dots, P-1\}$ and the vector of
 059 all time–frequency entries $\Phi\mathbf{x}(p, t, \lambda)$ for each path $p \in \mathcal{P}$ by $\phi_p(\mathbf{x})$, the ST loss function writes as:
 060

$$\mathcal{L}_{\mathbf{x}}^{\Phi}(\tilde{\mathbf{x}}) = \frac{1}{P} \sum_{p=0}^{P-1} \mathcal{L}_{\mathbf{x}}^{\phi_p}(\tilde{\mathbf{x}}) \quad \text{where} \quad \forall p \in \mathcal{P}, \mathcal{L}_{\mathbf{x}}^{\phi_p}(\tilde{\mathbf{x}}) = P \|\phi_p(\mathbf{x}) - \phi_p(\tilde{\mathbf{x}})\|^2. \quad (3)$$

061 Unfortunately, $\mathcal{L}_{\mathbf{x}}^{\Phi}(\tilde{\mathbf{x}})$ and its gradient $\nabla \mathcal{L}_{\mathbf{x}}^{\Phi}(\tilde{\mathbf{x}})$ are expensive in memory and in operations.
 062 Certainly, algorithmic refinements such as FFT-based filtering, multirate processing, and depth-first
 063 search can reduce the cost of an ST path (Oyallon et al., 2018). Yet, the need to traverse all P paths
 064 remains an obstacle to the applicability of multivariable ST for gradient-based learning at scale.

065 In this article, we aim to accelerate the training of an autoencoder F whose loss is ST distance
 066 between reference and reconstruction, and so over a finite corpus $\mathcal{X} = \{\mathbf{x}_0, \dots, \mathbf{x}_{N-1}\}$. Formally:

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{N} \sum_{n=0}^{N-1} (\mathcal{L}_{\mathbf{x}_n}^{\Phi} \circ F_{\mathbf{x}_n})(\mathbf{w}). \quad (4)$$

067 Given the decomposition in Equation 3, a naïve idea would be to replace each term $\mathcal{L}_{\mathbf{x}_n}^{\Phi}$ in the equation
 068 above by some per-path loss $\mathcal{L}_{\mathbf{x}_n}^{\phi_p}$, where the p ’s would be drawn independently and uniformly at
 069 random in the path set \mathcal{P} . This is a crude form of stochastic approximation (Benveniste et al.,
 070 2012) which is motivated by the tree-like structure of ST: neglecting the overhead of the first layer
 071 ($|\mathbf{W}\mathbf{x}|$), the computation of single-path gradient $\nabla \mathcal{L}_{\mathbf{x}_n}^{\phi_p}$ is roughly P times more efficient than that of
 072 a full ST gradient $\nabla \mathcal{L}_{\mathbf{x}_n}^{\Phi}$. However, this speedup comes at the detriment of numerical precision: a
 073 deterministic quantity has been replaced by an estimator whose variance may be impractically large.

074 “Scattering transform with Random Paths for machine Learning” (SCRAPL) is our proposed solution
 075 to this problem. Acknowledging that each single-path gradient makes for an inexpensive but noisy
 076 learning signal, we stabilize it via a combination of three stochastic optimization techniques and
 077 apply an architecture-informed importance sampling heuristic. Our contributions are:

- 078 1. **Stochastic approximation of scattering transform** through uniform sampling of paths.
- 079 2. **Path-wise adaptive moment estimation** (\mathcal{P} -Adam for short): an extension of the Adam
 algorithm (Kingma & Ba, 2014) which accounts for the non-i.i.d. nature of ST paths.
- 080 3. **Path-wise stochastic average gradient with acceleration** (\mathcal{P} -SAGA for short): a variant
 of the SAGA algorithm (Defazio et al., 2014) which keeps a memory of previous gradient
 values across all paths p .
- 081 4. **θ -importance sampling**: a parallelizable initialization heuristic that supplies auxiliary
 information to the stochastic optimizer by sampling paths p in proportion to the typical rate
 of change of the gradient in the optimization landscape.

082 Our main empirical finding is that SCRAPL accomplishes a favorable tradeoff between goodness of
 083 fit and computational efficiency on unsupervised sound matching, i.e., a nonlinear inverse problem in
 084 which the forward operator implements an audio synthesizer. In the context of differentiable digital
 085 signal processing (DDSP), the state-of-the-art perceptual loss function for this task is multiscale
 086 spectral loss (MSS, Yamamoto et al. (2020), Engel et al. (2020)). However, the gradient of MSS is
 087 uninformative when input and reconstruction are misaligned or when the synthesizer controls involve
 088 spectrotemporal modulations (Vahidi et al., 2023). Taking advantage of the stability guarantees of
 089 JTFS, SCRAPL expands the class of synthesizers which can be effectively decoded via DDSP.

090 Figure 1 illustrates one of our experiments: unsupervised sound matching for a nondeterministic
 091 granular synthesizer. On one hand, models based on MSS and other state-of-the-art perceptual losses
 092 are computationally efficient but inaccurate. On the other hand, JTFS-based models are five times
 093 more accurate but twenty five times more costly. SCRAPL is a new point on this Pareto front: it is
 094 within a factor two of JTFS in terms of accuracy while being within a factor two of MSS in terms of
 095 runtime, making it suitable for large-scale DDSP. Relatedly, SCRAPL is also more memory-efficient
 096 than JTFS, thus reducing overhead between CPU/GPU cores and allowing for a larger batch size.

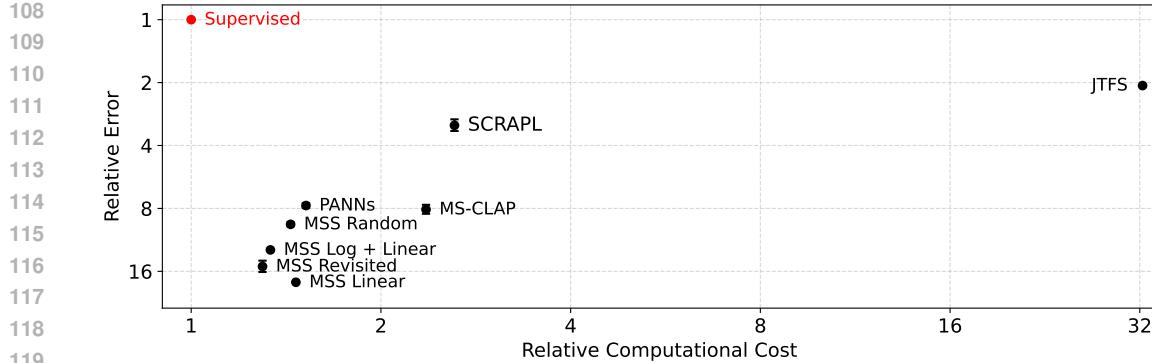


Figure 1: Mean average synthesizer parameter error (y-axis) versus computational cost (x-axis) of unsupervised sound matching models for the granular synthesis task. Both axes are rescaled by the performance of a supervised model with the same number of parameters. Whiskers denote 95% CI, estimated over 20 random seeds. Due to computational limitations, JTFS-based sound matching is evaluated only once.

2 RELATED WORK

The guiding intuition behind SCRAPL is that natural signals and images exhibit strong correlations across ST paths. This fact has been observed empirically since the onset of ST research (Bruna & Mallat, 2013; Andén & Mallat, 2011) and aligns with earlier work on texture modeling based on pairwise correlations between wavelet modulus coefficients (Portilla & Simoncelli, 2000).

Visual and auditory textures, understood as stationary random fields, play a key role in applied ST research. ST features outperform short-term Fourier features (e.g., MSS) in their ability to characterize intermittency in non-Gaussian textures (Muzy et al., 2015). Texture resynthesis by gradient descent of ST loss has been applied to such diverse settings as computer music creation (Lostanlen et al., 2019) and the study of the cosmic microwave background (Delouis et al., 2022).

The democratization of differentiable programming toolkits (e.g., TensorFlow, PyTorch, JAX) has greatly advanced the flexibility of gradient backpropagation in “hybrid” scattering–neural networks involving learnable and non-learnable modules. Angles & Mallat (2018) have built a hybrid scattering–GAN model for image generation, in which ST distance plays the role of a discriminator.

To our knowledge, the closest prior work to SCRAPL is the pruned graph scattering transform (pGST) of Ioannidis et al. (2020), a method which reduces the complexity of ST by pruning the path set \mathcal{P} down to a proper subset $\mathcal{P}' \subset \mathcal{P}$, based on a graph-spectrum-inspired criterion. Although both pGST and SCRAPL share a similar overarching goal, let us point out that pGST is a feature selection method: the cardinality of \mathcal{P}' is typically $\sim 10\%$ that of \mathcal{P} and \mathcal{P}' is kept fixed across training examples and across epochs. In comparison, SCRAPL performs a more radical pruning, down to a single path ($\text{card } \mathcal{P}' = 1$), while harnessing dedicated techniques in stochastic optimization (\mathcal{P} -Adam and \mathcal{P} -SAGA) to reduce the variance of ST loss during gradient backpropagation.

3 METHODS

3.1 STOCHASTIC APPROXIMATION OF SCATTERING TRANSFORM LOSS GRADIENT

The proposition below, proven in Appendix A, shows that if paths are drawn uniformly at random, the stochastic approximation in SCRAPL is unbiased: in other words, the expected value of the stochastic gradient of per-path loss is equal to the gradient of full ST loss.

Proposition 3.1. Let $\Phi = (\phi_p)_0^{P-1}$ be a scattering transform with P paths. Given a signal or image \mathbf{x} , let $F_{\mathbf{x}}$ be an autoencoder operating on \mathbf{x} and let $\mathcal{L}_{\mathbf{x}}^{\Phi}$ be the associated ST reconstruction loss. Let \mathcal{U}_P be the uniform distribution over $\mathcal{P} = \{0, \dots, P-1\}$. One has, for every weight vector \mathbf{w} :

$$\mathbb{E}_{z \sim \mathcal{U}_P} [\nabla (\mathcal{L}_{\mathbf{x}}^{\Phi} \circ F_{\mathbf{x}})(\mathbf{w})] = \nabla (\mathcal{L}_{\mathbf{x}}^{\Phi} \circ F_{\mathbf{x}})(\mathbf{w}). \quad (5)$$

162 Although a uniform sampling of paths matches the intuition of approximating the ST gradient in
 163 expectation, we will see that this may be suboptimal. The θ -importance sampling method, which
 164 we will present in Section 3.4, does not satisfy the hypothesis of Proposition 3.1; yet, it consistently
 165 outperforms uniform sampling as part of SCRAPL. The design of biased stochastic approximation
 166 schemes is an active topic in machine learning research (Dieuleveut et al., 2023).
 167

168 3.2 \mathcal{P} -ADAM: PATH-WISE ADAPTIVE MOMENT ESTIMATION

170 The key idea behind the Adam optimizer is to smooth the successive realizations of the stochastic
 171 gradient, here denoted by \mathbf{g} , via autoregressive estimates of its first- and second-order element-wise
 172 moments, denoted by \mathbf{m} and \mathbf{v} (Kingma & Ba, 2014). However, the smoothing technique in Adam is
 173 ineffective for SCRAPL because the gradients of path-wise ST losses are not identically distributed.
 174 Against this problem, we propose to maintain P estimates of path-wise moments (\mathcal{P} -Adam):

$$\mathbf{m}_p \leftarrow \beta_1^{(k-\tau_p)/P} \mathbf{m}_p + (1 - \beta_1^{(k-\tau_p)/P}) \mathbf{g} \quad (6)$$

$$\mathbf{v}_p \leftarrow \beta_2^{(k-\tau_p)/P} \mathbf{v}_p + (1 - \beta_2^{(k-\tau_p)/P}) (\mathbf{g} \odot \mathbf{g}), \quad (7)$$

175 where k is the current iteration number, τ_p is the iteration when path p was last drawn; β_1 and β_2 are
 176 hyperparameters; and the circled dot denotes element-wise multiplication of vectors. The exponent
 177 $(k - \tau_p)/P$ adapts the time constant of smoothing to the recency of the previous estimate.
 178

179 The second step in \mathcal{P} -Adam, following classical Adam, consists in bias correction and ratio of
 180 debiased first-order moment to stable square root of debiased second-order moment:
 181

$$\mathbf{g}_{\text{current}} = \frac{\frac{\mathbf{m}_p}{1 - \beta_1^{k/P}}}{\sqrt{\epsilon + \frac{\mathbf{v}_p}{1 - \beta_2^{k/P}}}}, \quad (8)$$

182 where we have adapted the original exponents of Adam (β_1^k, β_2^k) to account for the number of paths.
 183

184 3.3 \mathcal{P} -SAGA: PATH-WISE STOCHASTIC AVERAGE GRADIENT WITH ACCELERATION

185 The stochastic average gradient (SAG) algorithm has the potential to accelerate stochastic gradient
 186 descent in the context of the minimization of finite sums (Schmidt et al., 2017). Although this sum is
 187 typically over training examples in neural network training, in SCRAPL, Equation 3 is a sum over
 188 paths for a given example \mathbf{x} . With this observation in mind, we propose \mathcal{P} -SAGA, a path-wise
 189 version of SAG with acceleration (SAGA, Defazio et al. (2014)). We maintain a memory of the
 190 last \mathcal{P} -Adam updates over each path, denoted by $(\hat{\mathbf{g}}_p)_0^{P-1}$; and the set of paths previously visited,
 191 denoted by Γ . Given a learning rate α_k at iteration k , the \mathcal{P} -SAGA update is:
 192

$$\mathbf{w} \leftarrow \mathbf{w} - \alpha_k \left(\mathbf{g}_{\text{current}} - \hat{\mathbf{g}}_p + \frac{\sum_{\gamma \in \Gamma} \hat{\mathbf{g}}_\gamma}{\max(1, \text{card } \Gamma)} \right). \quad (9)$$

200 Unlike the original SAG and SAGA algorithms, \mathcal{P} -SAGA’s additional memory footprint is propor-
 201 tional to P , not the size of the dataset N , making it suitable for neural network training and real-world
 202 optimization tasks like our experiments in Section 4. We also note that \mathcal{P} -Adam and \mathcal{P} -SAGA
 203 introduce no additional hyperparameters over the standard Adam optimizer. Algorithm 1 summarizes
 204 SCRAPL with both \mathcal{P} -Adam and \mathcal{P} -SAGA enabled.
 205

206 3.4 θ -IMPORTANCE SAMPLING

207 We now consider the important special case of differentiable digital signal processing (DDSP, see
 208 Section 1), in which the autoencoder composes a non-learnable decoder, typically a synthesizer,
 209 with a learned encoder: i.e., $F_{\mathbf{x}} = (D \circ E_{\mathbf{x}})$ with $E_{\mathbf{x}}(\mathbf{w}) = \tilde{\theta}$ and $D(\tilde{\theta}) = \tilde{\mathbf{x}}$ (Engel et al., 2020).
 210 We assume both D and $E_{\mathbf{x}}$ to be differentiable with respect to their inputs, but D is not necessarily
 211 deterministic. We denote by U the dimension of the parameter space θ ; i.e., the output space of $E_{\mathbf{x}}$
 212 and input space of D .
 213

Algorithm 1 “Scattering transform with Random Paths for machine Learning” (SCRAPL). The pseudo-code below describes SCRAPL with a batch size equal to one, without loss of generality.

```

218 Require:  $\Phi = (\phi_p)_0^{P-1}$ : Scattering transform (ST) with  $P$  paths
219 Require:  $\pi$ : Categorical distribution over the path set  $\mathcal{P} = \{0, \dots, P-1\}$ 
220 Require:  $F$ : Autoencoder with trainable parameters  $w$ 
221 Require:  $w$ : Neural network weights at initialization
222 Require:  $\beta_1, \beta_2, \varepsilon$ : Adam hyperparameters
223 Require:  $(\alpha_k)_1^K$ : Learning rate schedule
224  $\Gamma \leftarrow \emptyset$ 
225 for  $p$  in  $\{0, \dots, P-1\}$  do
226    $\tau_p \leftarrow 0$ 
227    $m_p \leftarrow 0$ 
228    $v_p \leftarrow 0$ 
229    $\hat{g}_p \leftarrow 0$ 
230 end for
231 for  $k$  in  $\{1, \dots, K\}$  do
232    $n \leftarrow$  draw an integer uniformly at random in  $\{0, \dots, N-1\}$ 
233    $p \leftarrow$  draw an integer at random in  $\{0, \dots, P-1\}$  according to  $\pi$ 
234    $\mathcal{L}(w) \leftarrow P \|\phi_p(x_n) - (\phi_p \circ F_w)(x_n)\|_2^2$  {Stochastic approx.}
235    $\mathbf{g} \leftarrow \nabla \mathcal{L}(w)$ 
236    $m_p \leftarrow \beta_1^{(k-\tau_p)/P} m_p + (1 - \beta_1^{(k-\tau_p)/P}) \mathbf{g}$ 
237    $v_p \leftarrow \beta_2^{(k-\tau_p)/P} v_p + (1 - \beta_2^{(k-\tau_p)/P})(\mathbf{g} \odot \mathbf{g})$ 
238    $\hat{m} \leftarrow m_p / (1 - \beta_1^{k/P})$  { $\mathcal{P}$ -Adam}
239    $\hat{v} \leftarrow v_p / (1 - \beta_2^{k/P})$ 
240    $\mathbf{g}_{\text{current}} \leftarrow \hat{m} / \sqrt{\varepsilon + \hat{v}}$ 
241    $\tau_p \leftarrow k$ 
242
243    $\mathbf{g}_{\text{avg}} \leftarrow \frac{1}{\max(1, \text{card } \Gamma)} \sum_{\gamma \in \Gamma} \hat{g}_\gamma$ 
244    $\mathbf{g}_{\text{SAGA}} \leftarrow \mathbf{g}_{\text{current}} - \hat{g}_p + \mathbf{g}_{\text{avg}}$  { $\mathcal{P}$ -SAGA}
245    $w \leftarrow w - \alpha_k \mathbf{g}_{\text{SAGA}}$ 
246    $\hat{g}_p \leftarrow \mathbf{g}_{\text{current}}$ 
247    $\Gamma \leftarrow \Gamma \cup \{p\}$ 
248
249 end for
250 return  $w$ 
251

```

253 A known drawback of DDSP is that the optimization landscape of spectral loss in parameter space
 254 (i.e., of $\mathcal{L}_x^\Phi \circ D$) may not coincide with that of supervised parameter loss (i.e., Euclidean distance to
 255 θ , also known as P-loss) (Hayes et al., 2024). Against this drawback, we propose a method named
 256 θ -importance sampling (θ -IS), which constructs a categorical distribution π over the path space \mathcal{P} .
 257 The key idea behind θ -IS is to introduce bias in the stochastic approximation of spectral loss so as to
 258 bring it closer to P-loss. For lack of supervision, we are unable to construct the optimal distribution π
 259 but provide a heuristic of this form:

$$\pi_p = \frac{1}{U} \sum_{u=0}^{U-1} \frac{C_{u,p}}{\sum_{p=0}^{P-1} C_{u,p}}, \quad (10)$$

where, intuitively, $C_{u,p}$ represents the importance of parameter dimension θ_u upon path p . We rescale this importance relative to all paths and average uniformly across parameters u , yielding an importance-weighted categorical distribution π over paths. We then use π instead of a uniform distribution for sampling paths in the SCRAPL algorithm (see Algorithm 1).

Let $E_{\mathbf{x},u}(\mathbf{w})$ denote the u^{th} coordinate of $E_{\mathbf{x}}(\mathbf{w})$. Given \mathbf{w} , we measure the sensitivity of each ST path p to the parameter control u around the input \mathbf{x} in terms of the following partial derivative:

$$s_{\mathbf{x},u,p} : \mathbf{w} \longmapsto \frac{\partial (\mathcal{L}_{\mathbf{x}}^{\phi_p} \circ D)}{\partial \theta_c} \left(E_{\mathbf{x},u}(\mathbf{w}) \right) \quad (11)$$

To convert the sensitivity function $s_{\mathbf{x},u,p}$ into relative importance $C_{u,p}$, we analyze the curvature of the loss landscape $\mathcal{L}_{\mathbf{x}}^{\phi_p}$. We multiply $s_{\mathbf{x},u,p}$ by the gradient of $E_{\mathbf{x},u}$, yielding a linear transformation of neural network parameters. The coordinate-wise gradient of this linear transformation yields a positive definite matrix: we compute its largest eigenvalue. We repeat this process over a representative dataset \mathcal{X} of N_{IS} unlabeled signals from the training dataset. Formally:

$$C_{u,p} = \mathbb{E}_{\mathbf{x} \sim \mathcal{X}} [\lambda_{\max} (\nabla_{\mathbf{w}} (s_{\mathbf{x},u,p}(\mathbf{w}) \nabla E_{\mathbf{x},u}(\mathbf{w})))], \quad (12)$$

where $\lambda_{\max}(\mathbf{M})$ is the magnitude of the largest eigenvalue of a square matrix \mathbf{M} ; $\nabla_{\mathbf{w}}$ is the gradient with respect to \mathbf{w} . In practice, we compute $\lambda_{\max}(\mathbf{M})$ using a stochastic power iteration with deflation and the Hessian vector product, which has the same asymptotic runtime complexity as a backpropagation step.¹ Crucially, the computation required for θ -IS can be trivially parallelized across p and u , and only needs to be computed once before training.

Our definition of θ -IS is inspired by Schmidt et al. (2017), who propose a variant of the SAG algorithm in which mini-batches are sampled non-uniformly; more precisely, in proportion to the Lipschitz constant of the gradients, which we approximate using Equation 12 for each p and u . This heuristic relies on the argument that gradients which change quickly should be regarded as more important than gradients which change slowly.

4 EXPERIMENTS

We apply SCRAPL to a differentiable implementation of the joint time–frequency scattering transform (Muradeli et al., 2022). We conduct three unsupervised sound matching experiments under the DDSP autoencoder paradigm described in Section 3.4. The encoder, $E_{\mathbf{x}}$, is a convolutional neural network which operates on a constant- Q transform (Cheuk et al., 2020). We use relatively lightweight neural networks for our experiments, a choice made possible by the strong inductive bias DDSP provides and informed by prior work (Han et al., 2025) indicating that larger networks do not necessarily improve sound matching capabilities. We choose all hyperparameters in experiments heuristically.

To highlight the new kinds of perceptual quality assessment tasks SCRAPL enables, all three experiments investigate nondeterministic decoders that introduce random time shifts into the resulting reconstructed audio. While our experiments are for a discriminative and generative audio processing task, we emphasize that SCRAPL is a general algorithm and can be equally applied to deep inverse problems that leverage other scattering transforms.

4.1 JOINT TIME–FREQUENCY SCATTERING TRANSFORM (JTFS)

The joint time–frequency scattering transform (JTFS) is a nonlinear convolutional operator which extracts spectrotemporal modulations over the constant- Q spectrogram (Andén et al., 2019). The multivariable filter Ψ_p comprises two stages: temporal scattering, i.e., 1-D band-pass filtering with Morlet wavelets over the time axis; and frequential scattering, i.e., idem over the log-frequency axis. The center frequencies of band-pass filters for temporal scattering, called *rates*, are measured in Hertz. The center frequencies for frequential scattering, called *scales*, are measured in cycles per octave. Thus, in the case of JTFS, the path index p is a rate–scale multiindex.

The JTFS has been shown to correlate with human perception (Lostanlen et al., 2021; Tian et al., 2025) and can provide an informative gradient for audio comparisons that are misaligned (Vahidi et al., 2023) or benefit from multi-resolution analysis like percussive sounds (Han et al., 2024), which is why we select it as the underlying ST of the SCRAPL algorithm in our unsupervised sound matching experiments. Additionally, due to its computational complexity, until now it has been used almost exclusively as a precomputed feature instead of a loss function for neural network training.

4.2 GRANULAR SYNTH SOUND MATCHING

Granular synthesis is an example of a new class of synths that can be effectively sound matched with SCRAPL and the JTFS, due to its inherently stochastic audio generation process with individual grains being misaligned in time at the micro-level, but still being perceived as a single texture. It has been extensively used in the production of electronic music since the late 1950s² and played a

¹<https://github.com/noahgolmant/pytorch-hessian-eigenthings>

²<https://www.iannis-xenakis.org/en/granular-synthesis/>

fundamental role in the creation of contemporary electronic music genres. Our differentiable granular synth produces textures of chirplet grains with random temporal positions, center frequencies, and chirp rates, and has two continuous parameters: density (θ_{density}) which controls how many grains are produced, and slope (θ_{slope}) which controls their rate of frequency modulation.

We compare four MSS-based losses: linear, log + linear (Engel et al., 2020), random (Steinmetz & Reiss, 2020), and a SOTA hyperparameter-tuned revisited MSS loss (Schwär & Müller, 2023). Given their correlation with human perception (Kilgour et al., 2019; Tailleur et al., 2024), we also include the Euclidean distance of MS-CLAP (Elizalde et al., 2023) and PANNs Wavegram Logmel embeddings (Kong et al., 2020). In addition, we train with ordinary (i.e., full-tree) JTFS so as to put the speed and accuracy of SCRAPL into context. Lastly, as an estimate of best achievable performance with our encoder architecture and training configuration, we run a supervised version of sound matching, also known as “parameter loss” or P-loss for short (see Section 3.4). We summarize the implementation details in Appendix E.

4.3 CHIRPLET SYNTH SOUND MATCHING

Similar to the unsupervised granular synth sound matching experiment, we evaluate our θ -importance sampling initialization heuristic for SCRAPL on a differentiable chirplet synth (based on the implementation by Vahidi et al. (2023)) with two parameters: θ_{AM} which controls the rate of amplitude modulation (Hz) and θ_{FM} which controls the rate of frequency modulation (oct/s). Since the paths in the JTFS correspond to specific wavelet AM and FM center frequencies, given a chirplet synth configuration with bounded θ_{AM} and θ_{FM} ranges, we know approximately which paths of the JTFS should provide the most informative gradients for the synth parameters. After computing our initialization heuristic, we can analyze the resulting path probabilities and verify that the paths within the parameter ranges of the synth have been assigned a probability greater than uniform.

We evaluate four different synth configurations:

1. Slow AM ($\theta_{\text{AM}} \in [1.0, 2.0]$ Hz), slow FM ($\theta_{\text{FM}} \in [0.5, 1.0]$ oct/s);
2. Slow AM ($\theta_{\text{AM}} \in [1.0, 2.0]$ Hz), moderate FM ($\theta_{\text{FM}} \in [2.0, 4.0]$ oct/s);
3. Fast AM ($\theta_{\text{AM}} \in [2.8, 8.4]$ Hz), moderate FM ($\theta_{\text{FM}} \in [2.0, 4.0]$ oct/s);
4. Fast AM ($\theta_{\text{AM}} \in [2.8, 8.4]$ Hz), fast FM ($\theta_{\text{FM}} \in [4.0, 12.0]$ oct/s).

We compare SCRAPL training runs using uniform sampling and θ -importance sampling calculated from a single training batch of 32 examples. We summarize the implementation details in Appendix E.

4.4 ROLAND TR-808 SOUND MATCHING

As a real-world evaluation task, we sound match a DDSP implementation (Shier et al., 2024) of the Roland TR-808 Drum Machine, a historically meaningful synthesizer for the creation of Detroit techno, house, and hip-hop music.³ Inharmonic transient sounds like percussion are a form of non-stationary signal that the JTFS is well suited for perceptual quality assessment (Han et al., 2024) because of its ability to extract spectrotemporal patterns at multiple scales and rates. Additionally, due to the transient nature of drum sounds, they are highly sensitive to even a few milliseconds of misalignment, thus further benefiting from the time invariance of JTFS.

We use a high fidelity, 100% analog dataset⁴ of 681 bass drum, snare, tom, and hi-hat one-shot recordings of the TR-808 and repeat experiments 40 times on different train/validation/test splits and random seeds. Since the transient of analog drum recordings is rarely perfectly aligned, and no two analog TR-808 drum synths produce the same signal, we investigate both perfectly aligned (labeled *micro*) and unaligned (labeled *meso*) sound matching by up to ± 46 ms (± 2048 samples at 44.1 kHz). Given its correlation with human perception, we employ the JTFS and Fréchet Audio Distance (FAD) (Kilgour et al., 2019; Défossez et al., 2023) as evaluation metrics. We also include MSS and mean frame-by-frame perceptual loudness and loudness-weighted perceptually-scaled spectral centroid and flatness for both the transient and decay portions of reconstructed signals (nine metrics in total). Additional context for these last six metrics can be found in Shier et al. (2024). We summarize the implementation details in Appendix E.

³<https://www.ethanhein.com/wp/2016/beatmaking-fundamentals/>

⁴<https://samplesfrommars.com/products/tr-808-samples>

378
379
380
381
382

Table 1: Evaluation results for the unsupervised granular synth sound matching task with two continuous θ_{synth} parameters: θ_{density} and θ_{slope} (more details in Section 4.2). Uncertainties are 95% CI for 20 training runs using different random seeds. Due to computational limitations, the JTFS method is only evaluated once.

Method	$\theta_{\text{synth}} L_1 \text{ %o} \downarrow$	$\theta_{\text{density}} L_1 \text{ %o} \downarrow$	$\theta_{\text{slope}} L_1 \text{ %o} \downarrow$
JTFS	42.4	65.8	19.0
SCRAPL (no θ -IS)	73.8 \pm 13	70.4 \pm 8.8	77.2 \pm 19
SCRAPL	65.7 \pm 4.2	72.6 \pm 6.3	58.7 \pm 7.5
MSS Linear	370 \pm 0.52	499 \pm 0.84	241 \pm 0.28
MSS Log + Linear	259 \pm 1.7	277 \pm 3.2	241 \pm 0.42
MSS Revisited	311 \pm 19	376 \pm 40	246 \pm 3.0
MSS Random	195 \pm 4.2	149 \pm 7.8	242 \pm 1.0
MS-CLAP	166 \pm 8.2	81.9 \pm 9.0	250 \pm 8.2
PANNs Wavegram-Logmel	159 \pm 4.4	80.3 \pm 4.2	238 \pm 5.5
P-loss	20.5 \pm 0.20	24.7 \pm 0.31	16.3 \pm 0.31

5 RESULTS

5.1 GRANULAR SYNTH SOUND MATCHING

We benchmark all loss function computational costs (see Appendix B, Table 5) and plot them in Figure 1 against their evaluation accuracy (see Table 1) on $\theta_{\text{synth}} L_1$ relative to supervised training (i.e., P-loss). We observe that SCRAPL comes within factor two of JTFS in terms of accuracy, and within factor two of MSS in terms of runtime, striking a notable balance between them. The significant difference in runtime and convergence between JTFS and SCRAPL is further illustrated in Figure 2 where we plot validation accuracy against wall-clock time, instead of optimization steps. We also note that MSS is unable to sound match the synth at all, and the SOTA embedding losses are only able to optimize θ_{density} , albeit not as well as SCRAPL and JTFS. Validation accuracy curves for all methods are also provided in Figure 2. We provide a variation of Figure 1 plotting the test JTFS audio distance on the y-axis in Appendix B, Figure 3.

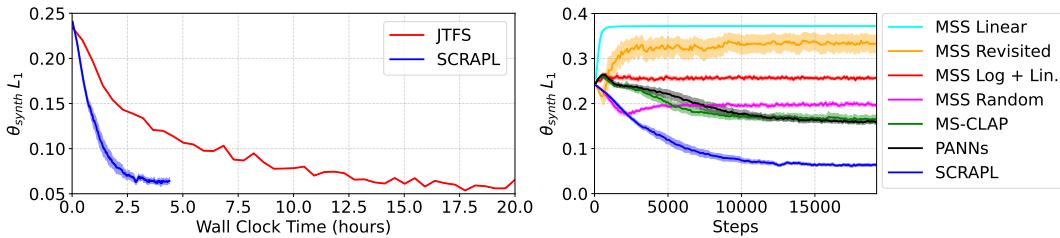


Figure 2: Left: JTFS vs. SCRAPL wall-clock training times on a single NVIDIA RTX A5000 GPU. Due to computational limitations, the JTFS method is only evaluated once. Right: Validation convergence graphs for the unsupervised granular synth sound matching task. Both: Shaded areas are 95% CI for 20 training runs using different random seeds.

Table 2 summarizes the results of an ablation of SCRAPL and its \mathcal{P} -Adam, \mathcal{P} -SAGA, and θ -IS optimization techniques for the granular synth sound matching task, with Appendix B, Table 6 providing additional information about statistical significance. There is a clear monotonic improvement in accuracy and convergence time for each technique, as well as a reduction in variance provided by \mathcal{P} -SAGA, and θ -IS. We emphasize that SCRAPL without any extra optimization techniques still outperforms all other non-JTFS methods in terms of accuracy, demonstrating its ability to optimize a new class of problems and making just stochastic sampling of scattering transforms a viable approach if the additional memory and computational requirements of \mathcal{P} -Adam, \mathcal{P} -SAGA, and θ -IS are undesirable. Finally, from Table 1, we see that θ -IS results in a better overall accuracy of θ_{synth} than uniform sampling (despite θ_{density} now being slightly worse), which is consistent with our hypothesis from Section 3.4 that θ -IS results in more balanced convergence of all synth parameters. Validation accuracy curves for all ablations are provided in Appendix B, Figure 4.

432
 433 Table 2: Ablation table for SCRAPL with test results and validation $\theta_{\text{synth}} L_1$ total variation and
 434 convergence steps for the unsupervised granular synth sound matching task. Convergence is defined
 435 as $\theta_{\text{synth}} L_1 < 100 \%$. Statistical significance results for each additional optimization technique are
 436 presented in Appendix B, Table 6. Uncertainties are 95% CI for 20 training runs using different
 437 random seeds. Due to computational limitations, the JTFS method is only evaluated once.

Method	\mathcal{P} -	\mathcal{P} -	θ -IS	Test		Validation	
	Adam	SAGA		$\theta_{\text{synth}} L_1 \downarrow \%$	Total Var. \downarrow	Conv. Steps \downarrow	
SCRAPL				99.7 \pm 8.2	5.30 \pm 0.25	10 906 \pm 1170	
	✓			87.4 \pm 15	6.98 \pm 0.25	8006 \pm 697	
	✓	✓		73.8 \pm 13	3.46 \pm 0.15	7296 \pm 683	
	✓	✓	✓	65.7 \pm 4.2	3.27 \pm 0.12	6014 \pm 642	
JTFS				42.4	5.66	1442	
P-loss				20.5 \pm 0.20	1.83 \pm 0.025	672 \pm 23	

448 In summary, this experiment demonstrates that the variance of the gradient elicited by the stochastic
 449 approximation of a ST with the SCRAPL algorithm is manageably small in the context of deep
 450 neural network (DNN) training, resulting in a favorable tradeoff between computational speed and
 451 convergence rate when compared to full-tree scattering (i.e., the JTFS). Training with SCRAPL is
 452 nearly equivalent to training with the gradient of full-tree scattering in terms of JTFS loss on unseen
 453 test data: see Appendix B, Figure 3. In terms of synthesizer parameter error, training with SCRAPL
 454 is not as accurate as training with full-tree ST, but outperforms prior work: see Table 1. This SOTA
 455 result paves the way towards a new kind of DNN training for deep inverse problems, in which the
 456 forward operator (synth) produces nondeterministic time–frequency patterns.

457 Understanding the convergence properties of algorithms like SCRAPL for convex and non-convex
 458 tasks is an active area of research (Reddi et al., 2016; 2018; Défossez et al., 2022; Kim & Oh, 2025).
 459 Our proof of Proposition 3.1 in Appendix A that SCRAPL without any additional optimization
 460 techniques is an unbiased estimator of full-tree ST is an important first step in this direction. We
 461 believe that further convergence analysis of SCRAPL remains a promising avenue for future work.

462 5.2 CHIRPLET SYNTH SOUND MATCHING

463 Table 3 summarizes the chirplet synth evaluation results, with Appendix C, Figure 5 showing
 464 validation accuracy curves for uniform and θ -importance sampling on the four synth configurations.
 465 θ -IS improves the prediction of θ_{AM} by 25–55% and of θ_{FM} by 14–80%, while reducing time to
 466 convergence by 23–50%: see Appendix C, Table 7. Of course, these improvements are for synth
 467 configurations that have been designed to showcase the benefit of nonuniform sampling of paths;
 468 however, this overall trend remains true, albeit not as pronounced, for the granular synth (Table 2) and
 469 real-world sound matching task (Table 4). Finally, we plot the path θ -IS probabilities in Appendix C,
 470 Figure 6 and observe that indeed, a unique distribution is learned for each synth, and the greater than
 471 uniform probabilities appear to roughly correspond to each configuration’s limited AM/FM range.

472 5.3 ROLAND TR-808 SOUND MATCHING

473 Table 4 and Appendix D, Tables 8, and 9 summarize the unsupervised Roland TR-808 synth sound
 474 matching audio distance, transient, and decay perceptual similarity results. Overall, we observe that
 475 JTFS dominates almost all metrics in both micro and meso environments, showcasing its suitability
 476 for transient percussive sounds and temporal invariance. After JTFS, MSS performs best when
 477 samples are perfectly aligned (micro), but performs worse in the unaligned (meso) setting and is
 478 unable to match the transient, which is the most salient part of a drum hit. In contrast, SCRAPL
 479 shows consistent sound matching performance in both micro and meso environments and is able to
 480 preserve the transient even when audio is misaligned. However, SCRAPL fails to recover the less
 481 audible decay portion of the signal. We hypothesize this is due to informative, low-frequency paths
 482 for the decay being sparse and underrepresented in the categorical distribution over paths, even after
 483 accounting for θ -IS. We provide listening samples at the accompanying website⁵ and encourage
 484 readers to evaluate the results directly.

485 ⁵Anonymous companion website: <https://icewithfrosting.github.io/scrapl/>

486
 487 Table 3: Evaluation results for SCRAPL with and without the θ -importance sampling initialization
 488 heuristic on unsupervised sound matching of four different AM / FM chirplet synths with two
 489 continuous θ_{synth} parameters: θ_{AM} and θ_{FM} (more details in Section 4.3). Uncertainties are 95% CI
 490 for 20 training runs using different random seeds.

Sampling Method (π)	Synth Configuration		θ_{AM}	$L_1 \text{ %o} \downarrow$	θ_{FM}	$L_1 \text{ %o} \downarrow$
	θ_{AM} (Hz)	θ_{FM} (oct/s)				
Uniform θ -IS	1.0 – 2.0	0.5 – 1.0	124 77.7	± 10 ± 6.7	155 78.4	± 18 ± 11
Uniform θ -IS	1.0 – 2.0	2.0 – 4.0	111 55.5	± 20 ± 4.1	68.6 44.4	± 11 ± 2.8
Uniform θ -IS	2.8 – 8.4	2.0 – 4.0	122 54.9	± 22 ± 3.5	238 48.5	± 21 ± 4.7
Uniform θ -IS	2.8 – 8.4	4.0 – 12.0	108 81.5	± 12 ± 12	95.6 82.1	± 20 ± 11

503
 504 Table 4: Audio distance evaluation results for the unsupervised Roland TR-808 DDSP synth sound
 505 matching task with 14 continuous θ_{synth} parameters (more details in Section 4.4). Uncertainties are
 506 95% CI for 40 training runs using different random seeds and dataset splits. Due to computational
 507 limitations, the JTFS method is only trained and evaluated for 4 random seeds.

Method	MSS Log. + Linear \downarrow		JTFS \downarrow		FAD (EnCodec) \downarrow	
	Micro	Meso	Micro	Meso	Micro	Meso
JTFS	617 ± 46	622 ± 45	490 ± 28	523 ± 17	0.781 ± 0.069	1.04 ± 0.15
SCRAPL (no θ -IS)	862 ± 36	944 ± 48	1140 ± 48	1250 ± 51	2.75 ± 0.39	2.85 ± 0.38
SCRAPL	857 ± 42	879 ± 42	1050 ± 50	1110 ± 52	2.43 ± 0.22	2.42 ± 0.22
MSS Lin.	611 ± 15	724 ± 37	779 ± 31	1470 ± 83	1.22 ± 0.082	3.33 ± 0.46
MSS L+L	596 ± 19	615 ± 18	1260 ± 58	1390 ± 49	2.14 ± 0.39	3.01 ± 0.40
MSS Rev.	637 ± 16	797 ± 20	870 ± 23	1250 ± 27	2.02 ± 0.37	2.21 ± 0.34
MSS Rand.	682 ± 25	700 ± 26	1410 ± 87	1500 ± 59	7.03 ± 2.2	6.65 ± 1.7

599 6 CONCLUSION

522 Differentiable similarity measures have the potential to enhance the perceptual quality of generative
 523 models and deep inverse problem solvers. In spite of their mathematical guarantees and neuro-
 524 physiological plausibility, scattering transforms (ST) have not been able to realize this potential,
 525 for lack of tractable optimization algorithms. To fill this gap, SCRAPL takes advantage of the
 526 tree-like structure of ST to save computation at each backward pass. Our numerical simulations
 527 show the value of SCRAPL for unsupervised sound matching, particularly when the synthesizer of
 528 interest is nondeterministic. Although our ST architecture of choice is joint time–frequency scattering
 529 (JTFS), we stress that SCRAPL is agnostic to the specifics of multivariable filterbank design: beyond
 530 wavelet scattering, it extends to learnable scattering-like architectures (Lattner et al., 2019; Cotter &
 531 Kingsbury, 2019; Gauthier et al., 2022). We consider investigating SCRAPL’s generalizability to
 532 other ST architectures, different audio tasks such as speech enhancement and automatic mixing, and
 533 additional modalities like adversarial image generation and texture synthesis as promising directions
 534 for future work. As a longer-term perspective, the success of our architecture-informed importance
 535 sampling heuristic highlights the opportunity to meta-learn the relative importance of each ST path
 536 for the task at hand over the course of neural network training (Yamaguchi et al., 2023).

540 REPRODUCIBILITY STATEMENT
541

542 Appendix E contains all hyperparameters and training details for each of the three experiments in
 543 this paper. We also provide listening samples, anonymized source code, configuration files, and
 544 instructions to reproduce our experiments at the anonymous companion website:
 545 <https://icewithfrosting.github.io/scrap/>

546
547 REFERENCES

- 548 Joakim Andén and Stéphane Mallat. Multiscale scattering for audio classification. In *Proceedings of*
 549 *the International Society for Music Information Retrieval Conference (ISMIR)*. Miami, Florida,
 550 2011.
- 552 Joakim Andén and Stéphane Mallat. Deep scattering spectrum. *IEEE Transactions on Signal*
 553 *Processing*, 62(16):4114–4128, 2014.
- 554 Joakim Andén, Vincent Lostanlen, and Stéphane Mallat. Joint time–frequency scattering. *IEEE*
 555 *Transactions on Signal Processing*, 67(14):3704–3718, 2019.
- 557 Tomás Angles and Stéphane Mallat. Generative networks as inverse problems with scattering
 558 transforms. In *Proceedings of the International Conference on Learning Representations (ICLR)*,
 559 2018.
- 560 Albert Benveniste, Michel Métivier, and Pierre Priouret. *Adaptive algorithms and stochastic approxi-*
 561 *mations*, volume 22. Springer Science & Business Media, 2012.
- 562 Joan Bruna and Stéphane Mallat. Invariant scattering convolution networks. *IEEE transactions on*
 563 *pattern analysis and machine intelligence*, 35(8):1872–1886, 2013.
- 565 Kin Wai Cheuk, Hans Anderson, Kat Agres, and Dorien Herremans. nnaudio: An on-the-fly gpu
 566 audio to spectrogram conversion toolbox using 1d convolutional neural networks. *IEEE Access*, 8:
 567 161981–162003, 2020. doi: 10.1109/ACCESS.2020.3019084.
- 568 Fergal Cotter and Nick Kingsbury. A learnable ScatterNet: Locally invariant convolutional layers. In
 569 *Proceedings of the IEEE International Conference on Image Processing (ICIP)*. IEEE, 2019.
- 571 Aaron Defazio, Francis Bach, and Simon Lacoste-Julien. SAGA: A fast incremental gradient method
 572 with support for non-strongly convex composite objectives. *Advances in Neural Information*
 573 *Processing Systems (NeurIPS)*, 2014.
- 574 Alexandre Défossez, Leon Bottou, Francis Bach, and Nicolas Usunier. A simple convergence proof
 575 of adam and adagrad. *Transactions on Machine Learning Research*, 2022. ISSN 2835-8856. URL
 576 <https://openreview.net/forum?id=ZPQhzTSWA7>.
- 577 Alexandre Défossez, Jade Copet, Gabriel Synnaeve, and Yossi Adi. High fidelity neural audio
 578 compression. *Transactions on Machine Learning Research*, 2023. ISSN 2835-8856. https:
 579 <https://openreview.net/forum?id=ivCd8z8zR2>.
- 581 J-M Delouis, Erwan Ally, Edouard Gauvrit, and François Boulanger. Non-gaussian modelling
 582 and statistical denoising of planck dust polarisation full-sky maps using scattering transforms.
 583 *Astronomy & Astrophysics*, 668:A122, 2022.
- 584 Aymeric Dieuleveut, Gersende Fort, Eric Moulines, and Hoi-To Wai. Stochastic approximation
 585 beyond gradient for signal processing and machine learning. *IEEE Transactions on Signal*
 586 *Processing*, 71:3117–3148, 2023.
- 587 Benjamin Elizalde, Soham Deshmukh, Mahmoud Al Ismail, and Huaming Wang. Clap learning
 588 audio concepts from natural language supervision. In *IEEE International Conference on Acoustics,*
 589 *Speech and Signal Processing (ICASSP)*, pp. 1–5, 2023. doi: 10.1109/ICASSP49357.2023.
 590 10095889.
- 592 Jesse Engel, Lamtharn Hantrakul, Chenjie Gu, and Adam Roberts. DDSP: Differentiable digital
 593 signal processing. *Proceedings of the International Conference on Learning Representations*
 (ICLR), 2020.

- 594 Shanel Gauthier, Benjamin Thérien, Laurent Alsene-Racicot, Muawiz Chaudhary, Irina Rish, Eugene
 595 Belilovsky, Michael Eickenberg, and Guy Wolf. Parametric scattering networks. In *Proceedings of*
 596 *the IEEE International Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022.
 597
- 598 Han Han, Vincent Lostanlen, and Mathieu Lagrange. Learning to solve inverse problems for
 599 perceptual sound matching. *IEEE/ACM Transactions on Audio, Speech, and Language Processing*,
 600 32:2605–2615, 2024.
- 601 Han Han, Vincent Lostanlen, and Mathieu Lagrange. Gradient Clipping Improves Neural Network
 602 Optimization for Perceptual Sound Matching. In *Proceedings of the European Signal Processing*
 603 *Conference (EUSIPCO)*, Palermo, Italy, September 2025. URL <https://hal.science/hal-05124224>.
- 604 Ben Hayes, Jordie Shier, György Fazekas, Andrew McPherson, and Charalampos Saitis. A review
 605 of differentiable digital signal processing for music and speech synthesis. *Frontiers in Signal*
 606 *Processing*, 3:1284100, 2024.
- 607 Vassilis N. Ioannidis, Siheng Chen, and Georgios B. Giannakis. Pruned graph scattering transforms.
 608 In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2020.
- 609 Kevin Kilgour, Mauricio Zuluaga, Dominik Roblek, and Matthew Sharifi. Fréchet audio distance:
 610 A reference-free metric for evaluating music enhancement algorithms. In *Proceedings of the*
 611 *International Speech Communication Association Conference (Interspeech)*, pp. 2350–2354, 2019.
 612 doi: 10.21437/Interspeech.2019-2219.
- 613 Gyu Yeol Kim and Min-hwan Oh. ADAM optimization with adaptive batch selection. In *The*
 614 *Thirteenth International Conference on Learning Representations*, 2025. URL <https://openreview.net/forum?id=BZrSCv2SBq>.
- 615 Diederik P Kingma and Jimmy Ba. Adam: A method for stochastic optimization. In *Proceedings of*
 616 *the International Conference on Learning Representations (ICLR)*, 2014.
- 617 Qiuqiang Kong, Yin Cao, Turab Iqbal, Yuxuan Wang, Wenwu Wang, and Mark D. Plumbley.
 618 Panks: Large-scale pretrained audio neural networks for audio pattern recognition. *IEEE/ACM*
 619 *Trans. Audio, Speech and Lang. Proc.*, 28:2880–2894, November 2020. ISSN 2329-9290. doi:
 620 10.1109/TASLP.2020.3030497. URL <https://doi.org/10.1109/TASLP.2020.3030497>.
- 621 Nina Kowalski, Didier A Depireux, and Shihab A Shamma. Analysis of dynamic spectra in ferret
 622 primary auditory cortex. i. characteristics of single-unit responses to moving ripple spectra. *Journal*
 623 *of Neurophysiology*, 76(5):3503–3523, 1996.
- 624 Stefan Lattner, Monika Dörfler, and Andreas Arzt. Learning complex basis functions for invariant
 625 representations of audio. In *Proceedings of the International Society Society for Music Information*
 626 *Retrieval Conference (ISMIR)*, 2019.
- 627 Vincent Lostanlen and Stéphane Mallat. Wavelet scattering on the pitch spiral. In *Proceedings of the*
 628 *Digital Audio Effects Conference (DAFx)*, 2016.
- 629 Vincent Lostanlen, Joakim Andén, and Mathieu Lagrange. Fourier at the heart of computer music:
 630 From harmonic sounds to texture. *Comptes Rendus. Physique*, 20(5):461–473, 2019.
- 631 Vincent Lostanlen, Christian El-Hajj, Mathias Rossignol, Grégoire Lafay, Joakim Andén, and Mathieu
 632 Lagrange. Time–frequency scattering accurately models auditory similarities between instrumental
 633 playing techniques. *EURASIP Journal on Audio, Speech, and Music Processing*, 2021(1):3, 2021.
- 634 Stéphane Mallat. Group invariant scattering. *Communications on Pure and Applied Mathematics*, 65
 635 (10):1331–1398, 2012.
- 636 Stéphane Mallat. Understanding deep convolutional networks. *Philosophical Transactions of the*
 637 *Royal Society A: Mathematical, Physical and Engineering Sciences*, 374(2065):20150203, 2016.
- 638 John Muradeli, Cyrus Vahidi, Changhong Wang, Han Han, Vincent Lostanlen, Mathieu Lagrange,
 639 and George Fazekas. Differentiable time-frequency scattering on gpu. In *Proceedings of the*
 640 *Digital Audio Effects Conference (DAFx)*, 2022.

- 648 Jean-François Muzy, Emmanuel Bacry, Stéphane Mallat, and Joan Bruna. Intermittent process
 649 analysis with scattering moments. *Annals of Statistics*, 43(1):323, 2015.
 650
- 651 Sam V Norman-Haignere and Josh H McDermott. Neural responses to natural and model-matched
 652 stimuli reveal distinct computations in primary and nonprimary auditory cortex. *PLoS biology*, 16
 653 (12):e2005127, 2018.
- 654 Edouard Oyallon, Stéphane Mallat, and Laurent Sifre. Generic deep networks with wavelet scattering.
 655 In *Proceedings of the International Conference on Learning Representations (ICLR) Workshop
 656 Track*, 2014.
- 657
- 658 Edouard Oyallon, Sergey Zagoruyko, Gabriel Huang, Nikos Komodakis, Simon Lacoste-Julien,
 659 Matthew Blaschko, and Eugene Belilovsky. Scattering networks for hybrid representation learning.
 660 *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 41(9):2208–2221, 2018.
- 661 Kailash Patil, Daniel Pressnitzer, Shihab Shamma, and Mounya Elhilali. Music in our ears: The
 662 biological bases of musical timbre perception. *PLoS computational biology*, 8(11):e1002759,
 663 2012.
- 664
- 665 Javier Portilla and Eero P Simoncelli. A parametric texture model based on joint statistics of complex
 666 wavelet coefficients. *International Journal of Computer Vision*, 40(1):49–70, 2000.
- 667
- 668 Sashank J. Reddi, Suvrit Sra, Barnabás Póczos, and Alex Smola. Fast incremental method for smooth
 669 nonconvex optimization. In *2016 IEEE 55th Conference on Decision and Control (CDC)*, pp.
 670 1971–1977. IEEE Press, 2016. doi: 10.1109/CDC.2016.7798553. URL <https://doi.org/10.1109/CDC.2016.7798553>.
- 671
- 672 Sashank J. Reddi, Satyen Kale, and Sanjiv Kumar. On the convergence of adam and beyond. In
 673 *International Conference on Learning Representations*, 2018. URL <https://openreview.net/forum?id=ryQu7f-RZ>.
- 674
- 675
- 676 Mark Schmidt, Nicolas Le Roux, and Francis Bach. Minimizing finite sums with the stochastic
 677 average gradient. *Math. Program.*, 162(1–2):83–112, March 2017. ISSN 0025-5610. doi:
 678 10.1007/s10107-016-1030-6. URL <https://doi.org/10.1007/s10107-016-1030-6>.
- 679
- 680 Simon Schwär and Meinard Müller. Multi-scale spectral loss revisited. *IEEE Signal Processing
 681 Letters*, 30:1712–1716, 2023. doi: 10.1109/LSP.2023.3333205.
- 682
- 683 Jordie Shier, Charalampos Saitis, Andrew Robertson, and Andrew McPherson. Real-time timbre
 684 remapping with differentiable dsp. In *Proceedings of the International Conference on New
 685 Interfaces for Musical Expression NIME*, 2024. URL <https://arxiv.org/abs/2407.04547>.
- 686
- 687 Laurent Sifre and Stéphane Mallat. Rotation, scaling and deformation invariant scattering for
 688 texture discrimination. In *Proceedings of the IEEE Conference on Computer Vision and Pattern
 689 Recognition*, pp. 1233–1240, 2013.
- 690
- 691 Christian J. Steinmetz and Joshua D. Reiss. auraloss: Audio focused loss functions in PyTorch. In
 692 *Digital Music Research Network One-day Workshop (DMRN+15)*, 2020.
- 693
- 694 Modan Tailleur, Junwon Lee, Mathieu Lagrange, Keunwoo Choi, Laurie M. Heller, Keisuke Imoto,
 695 and Yuki Okamoto. Correlation of fréchet audio distance with human perception of environmental
 696 audio is embedding dependant. In *Proceedings of the European Signal Processing Conference
 697 (EUSIPCO)*, 2024.
- 698
- 699 Haokun Tian, Stefan Lattner, and Charalampos Saitis. Assessing the alignment of audio representa-
 700 tions with timbre similarity ratings. *Proceedings of the International Society for Music Information
 701 Retrieval Conference (ISMIR)*, 2025.
- 702
- Cyrus Vahidi, Han Han, Changhong Wang, Mathieu Lagrange, György Fazekas, and Vincent
 703 Lostanlen. Mesostructures: Beyond spectrogram loss in differentiable time-frequency analy-
 704 sis. *Journal of the Audio Engineering Society*, 71(9):577–585, 2023.

- 702 Shin'ya Yamaguchi, Daiki Chijiwa, Sekitoshi Kanai, Atsutoshi Kumagai, and Hisashi Kashima. Reg-
703 ularizing neural networks with meta-learning generative models. *Advances in Neural Information*
704 *Processing Systems (NeurIPS)*, 36:27315–27331, 2023.
- 705
706 Ryuichi Yamamoto, Eunwoo Song, and Jae-Min Kim. Parallel wavegan: A fast waveform generation
707 model based on generative adversarial networks with multi-resolution spectrogram. *Proceedings*
708 *of IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2020.
- 709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755

756 A PROOF OF PROPOSITION 3.1
 757

758 Let us write $\tilde{\mathbf{x}} = F_{\mathbf{x}}(\mathbf{w})$. By linearity of the gradient, we may decompose $\nabla \mathcal{L}_{\mathbf{x}}^{\Phi}(\tilde{\mathbf{x}})$ over paths:
 759

$$760 \quad \nabla \mathcal{L}_{\mathbf{x}}^{\Phi}(\tilde{\mathbf{x}}) = \frac{1}{P} \sum_{p=0}^{P-1} \nabla \mathcal{L}_{\mathbf{x}}^{\phi_p}(\tilde{\mathbf{x}}). \quad (13)$$

761
762

763 Let us denote the Jacobian of $F_{\mathbf{x}}$ at \mathbf{w} by $\mathbf{J}_{F_{\mathbf{x}}}(\mathbf{w})$. For each $p \in \mathcal{P}$, we apply the chain rule:
 764

$$765 \quad \nabla (\mathcal{L}_{\mathbf{x}}^{\phi_p} \circ F_{\mathbf{x}})(\mathbf{w}) = \mathbf{J}_{F_{\mathbf{x}}}(\mathbf{w})^{\top} \nabla \mathcal{L}_{\mathbf{x}}^{\phi_p}(\tilde{\mathbf{x}}). \quad (14)$$

766
767

768 Plugging the identity above into Equation 13 yields:
 769

$$770 \quad \begin{aligned} \nabla (\mathcal{L}_{\mathbf{x}}^{\Phi} \circ F_{\mathbf{x}})(\mathbf{w}) &= \frac{1}{P} \sum_{p=0}^{P-1} \left(\mathbf{J}_{F_{\mathbf{x}}}(\mathbf{w})^{\top} \nabla \mathcal{L}_{\mathbf{x}}^{\phi_p}(\tilde{\mathbf{x}}) \right) \\ 771 &= \mathbf{J}_{F_{\mathbf{x}}}(\mathbf{w})^{\top} \left(\frac{1}{P} \sum_{p=0}^{P-1} \nabla \mathcal{L}_{\mathbf{x}}^{\phi_p}(\tilde{\mathbf{x}}) \right), \end{aligned} \quad (15)$$

772
773

774 where the latter equation holds by associativity of matrix multiplication.
 775

776 We now compute the expected value of $\nabla \mathcal{L}_{\mathbf{x}}^{\phi_z}(\tilde{\mathbf{x}})$ for $z \sim \mathcal{U}_P$, i.e., a uniform distribution over \mathcal{P} :
 777

$$778 \quad \mathbb{E}_{z \sim \mathcal{U}_P} [\nabla \mathcal{L}_{\mathbf{x}}^{\phi_z}(\tilde{\mathbf{x}})] = \frac{1}{P} \sum_{p=0}^{P-1} \nabla \mathcal{L}_{\mathbf{x}}^{\phi_p}(\tilde{\mathbf{x}}). \quad (16)$$

779

780 We recognize the column vector on the right-hand side of Equation 15. Thus:
 781

$$782 \quad \begin{aligned} \nabla (\mathcal{L}_{\mathbf{x}}^{\Phi} \circ F_{\mathbf{x}})(\mathbf{w}) &= \mathbf{J}_{F_{\mathbf{x}}}(\mathbf{w})^{\top} \mathbb{E}_{z \sim \mathcal{U}_P} [\nabla \mathcal{L}_{\mathbf{x}}^{\phi_z}(\tilde{\mathbf{x}})] \\ 783 &= \mathbb{E}_{z \sim \mathcal{U}_P} [\mathbf{J}_{F_{\mathbf{x}}}(\mathbf{w})^{\top} \nabla \mathcal{L}_{\mathbf{x}}^{\phi_z}(\tilde{\mathbf{x}})], \end{aligned} \quad (17)$$

784

785 where the latter equation holds by linearity of the expected value. Finally, we use the reverse form of
 786 the chain rule in Equation 14 to identify the expected SCRAPL gradient:
 787

$$788 \quad \nabla (\mathcal{L}_{\mathbf{x}}^{\Phi} \circ F_{\mathbf{x}})(\mathbf{w}) = \mathbb{E}_{z \sim \mathcal{U}_P} [\nabla (\mathcal{L}_{\mathbf{x}}^{\phi_z} \circ F_{\mathbf{x}})(\mathbf{w})] \quad (18)$$

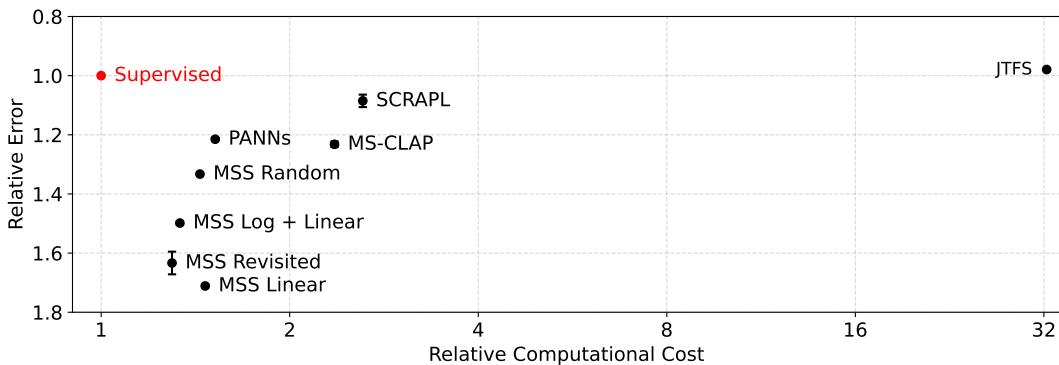
789 concluding the proof.

790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809

810 **B ADDITIONAL GRANULAR SYNTH EVALUATION RESULTS**
 811
 812

813 Table 5: Loss function benchmark results for one optimization step (forward + backward, 32768
 814 samples of audio, batch size 4, 1 thread, single precision, 1 NVIDIA RTX A5000 GPU, CUDA 12.4,
 815 PyTorch 2.8.0). SCRAPL paths are benchmarked individually and then aggregated across all paths
 816 using the median for time and interquartile range (IQR), and maximum for memory usage.

Method	Median Time (ms) ↓	IQR (ms) ↓	Max. Memory (MB) ↓
JTFS	1730	23.9	12 967
SCRAPL	89.8	3.62	2503
MSS Linear	26.3	1.12	694
MSS Log + Linear	19.1	0.696	702
MSS Revisited	17.0	0.210	663
MSS Random	24.7	5.81	706
MS-CLAP	75.6	1.69	2032
PANNs Wavegram-Logmel	29.3	5.92	1360
P-loss ($\dim(\theta_{\text{synth}}) = 2$)	0.516	0.108	625



811 Figure 3: Mean average JTFS perceptual audio distance (y-axis) versus computational cost (x-axis)
 812 of unsupervised sound matching models for the granular synthesis task. Both axes are rescaled by the
 813 performance of a supervised model with the same number of parameters. Whiskers denote 95% CI,
 814 estimated over 20 random seeds. Due to computational limitations, JTFS-based sound matching is
 815 evaluated only once.

816
 817
 818 Table 6: Statistical significance and relative improvement of each additional SCRAPL optimization
 819 technique for the unsupervised granular synth sound matching task. Convergence is defined as
 820 $\theta_{\text{synth}} L_1 < 100 \%$. See Table 2 for absolute results and comparisons to JTFS and P-loss. Uncertainties
 821 are 95% CI for 20 training runs using different random seeds. Due to computational limitations, the
 822 JTFS method is only evaluated once.

Opt.	Test			Validation		
	Technique	$\Delta \theta_{\text{synth}} L_1 \% \downarrow$	$\Delta \text{Total Var.} \downarrow$		$\Delta \text{Conv. Steps} \downarrow$	
+ \mathcal{P} -Adam		$-12.3 \pm 17 \quad (p = 0.15)$	$1.68 \pm 0.35 \quad (p < 0.01)$	$-2900 \pm 1800 \quad (p < 0.01)$		
+ \mathcal{P} -SAGA		$-13.6 \pm 5.6 \quad (p < 0.01)$	$-3.52 \pm 0.18 \quad (p < 0.01)$	$-710 \pm 450 \quad (p < 0.01)$		
+ θ -IS		$-8.13 \pm 15 \quad (p = 0.26)$	$-0.188 \pm 0.092 \quad (p < 0.01)$	$-1280 \pm 410 \quad (p < 0.01)$		

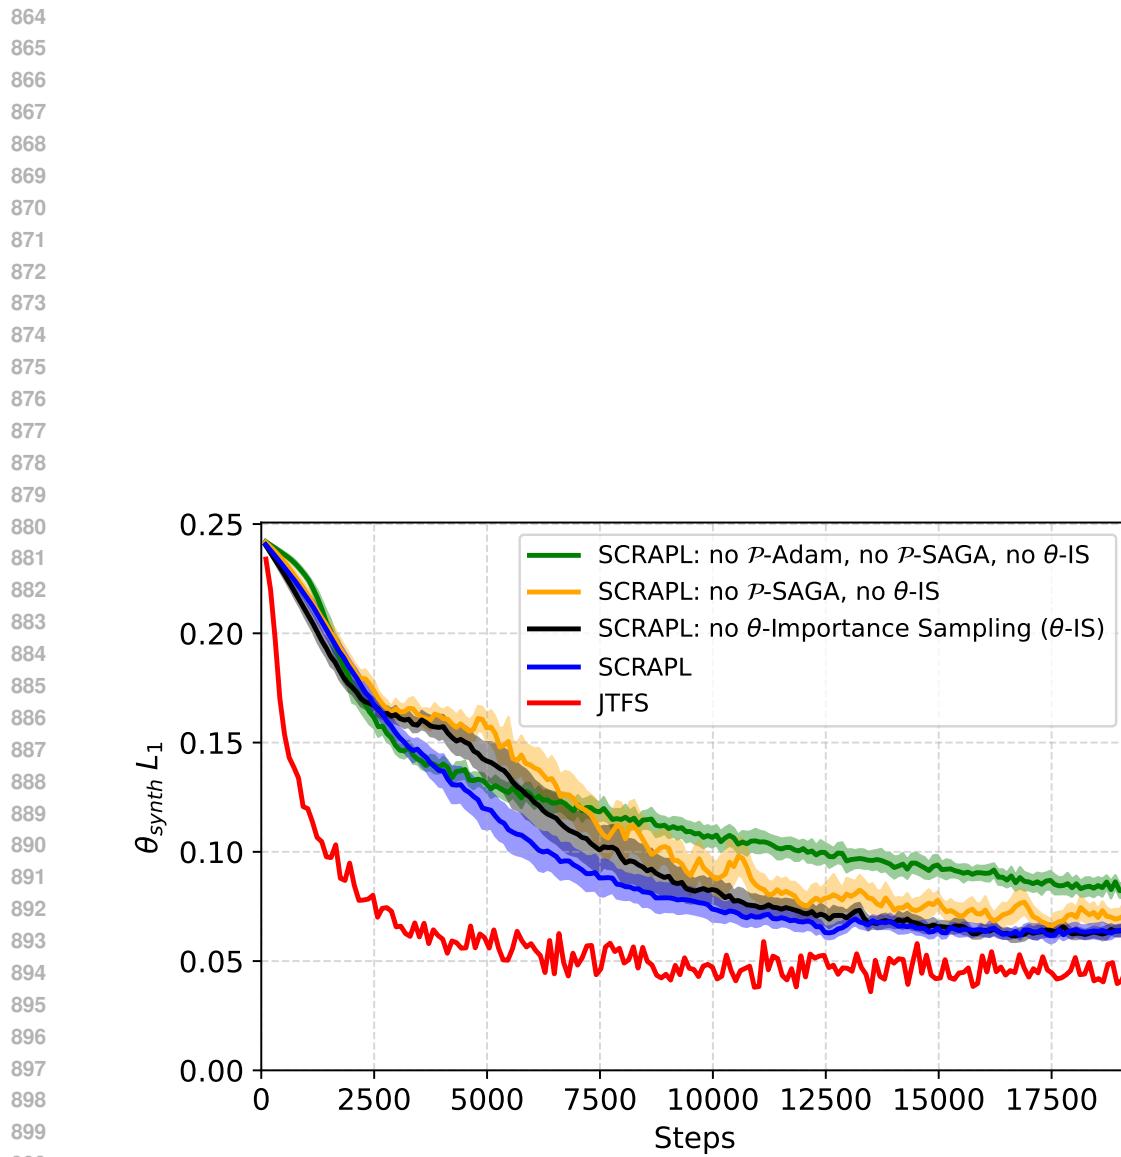
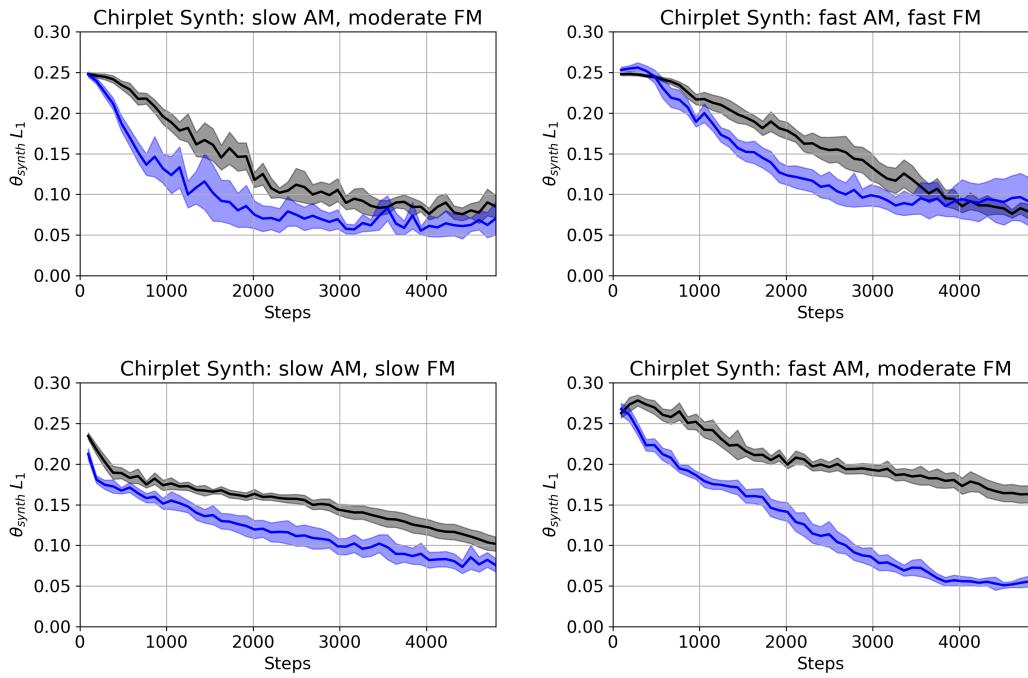


Figure 4: Validation convergence graphs of SCRAPL ablations and the JTFS for the unsupervised granular synth sound matching task. Shaded areas are 95% CI for 20 training runs using different random seeds. Due to computational limitations, the JTFS method is only evaluated once.

918
 919
 920
C ADDITIONAL CHIRPLET SYNTH EVALUATION RESULTS
 921
 922
 923
 924
 925
 926
 927
 928
 929
 930
 931



922
 923
 924
 925
 926
 927
 928
 929
 930
 931
 932
 933
 934
 935
 936
 937
 938
 939
 940
 941
 942
 943
 944
 945
 946
 947
 948
 949
 950
 951
 952
 953
 954
 955
 956
 957
 958
 959
 960
 961
 962
 963
 964
 965
 966
 967
 968
 969
 970
 971

Figure 5: SCRAPL θ_{synth} L_1 validation values during training for four different AM / FM chirplet synths with two continuous θ_{synth} parameters: θ_{AM} and θ_{FM} (more details in Section 4.3). Blue is using the θ -importance sampling initialization heuristic, and black is using uniform sampling. Shaded areas are 95% CI for 20 training runs using different random seeds.

Table 7: Convergence rate (CR) and steps for SCRAPL with and without the θ -importance sampling initialization heuristic on unsupervised sound matching of four different AM / FM chirplet synths with two continuous θ_{synth} parameters: θ_{AM} and θ_{FM} (more details in Section 4.3). Convergence is defined as $L_1 < 100\%$ for θ_{AM} or θ_{FM} . Uncertainties are 95% CI for 20 training runs using different random seeds.

Sampling Method (π)	Synth Configuration		θ_{AM}		θ_{FM}	
	θ_{AM} (Hz)	θ_{FM} (oct/s)	CR \uparrow	Conv. Steps \downarrow	CR \uparrow	Conv. Steps \downarrow
Uniform θ -IS	1.0 – 2.0	0.5 – 1.0	60% 100%	3944 ± 342 2002 ± 324	45% 100%	4064 ± 372 3134 ± 492
Uniform θ -IS	1.0 – 2.0	2.0 – 4.0	100% 100%	2203 ± 135 1099 ± 173	100% 100%	1536 ± 194 768 ± 118
Uniform θ -IS	2.8 – 8.4	2.0 – 4.0	95% 100%	3254 ± 250 1925 ± 165	0% 100%	N/A 2966 ± 210
Uniform θ -IS	2.8 – 8.4	4.0 – 12.0	100% 95%	3096 ± 334 2253 ± 218	95% 95%	3208 ± 235 2178 ± 173

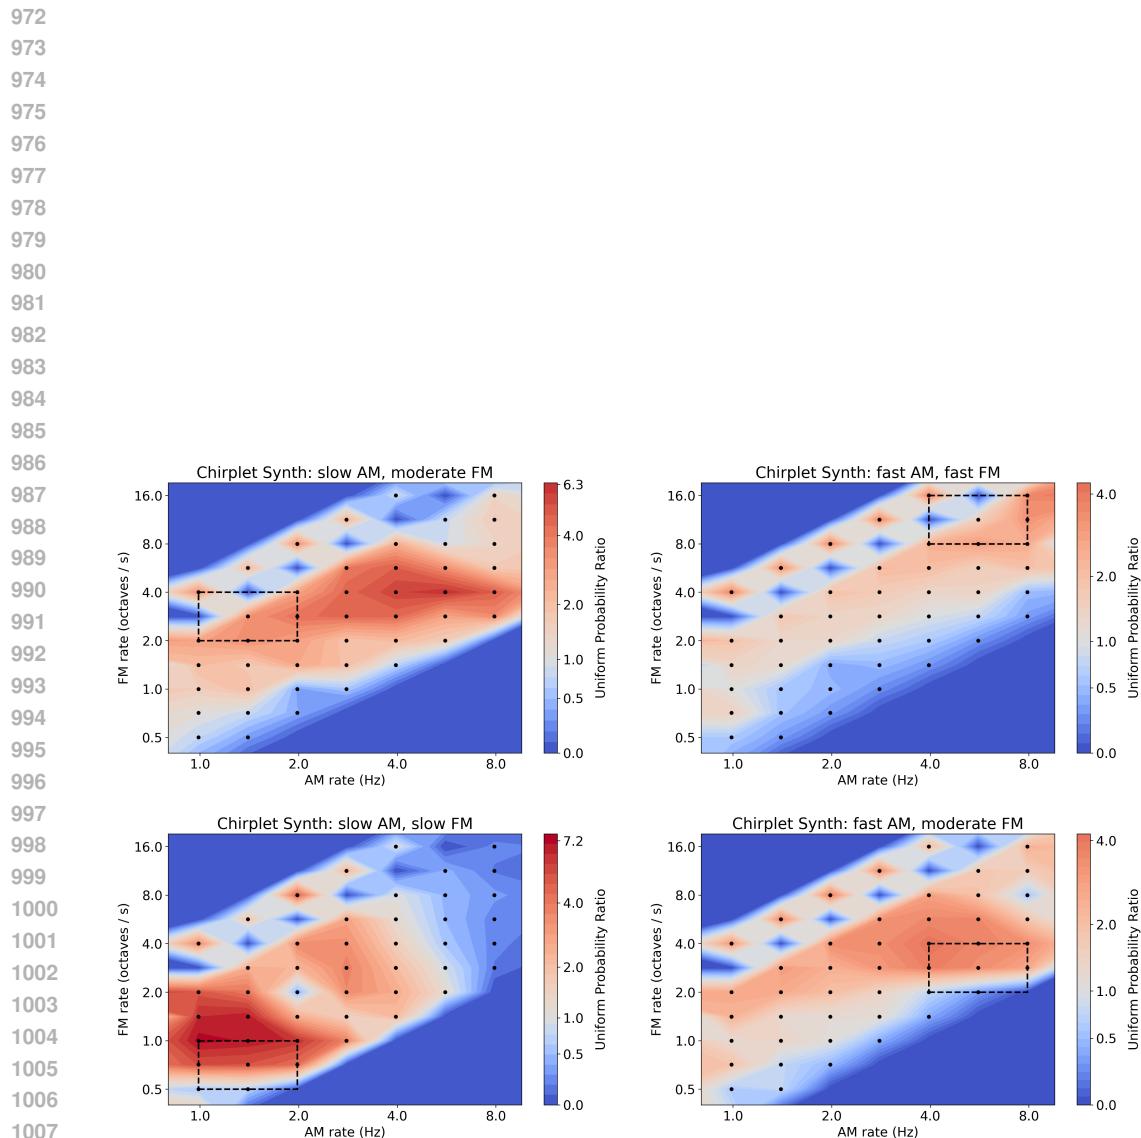


Figure 6: SCRABL path θ -importance sampling probabilities for four different AM / FM chirplet synths calculated from 1 batch of 32 log-uniformly randomly sampled θ_{synth} values. Black dots are individual path (wavelet) AM / FM center frequency locations and each dashed rectangle is the θ_{synth} range for each synth configuration. A uniform probability ratio of 1.0 means a path is sampled with probability $1/P$.

1026 **D ADDITIONAL ROLAND TR-808 EVALUATION RESULTS**
 1027
 1028

1029 Table 8: Drum transient evaluation results for the unsupervised Roland TR-808 DDSP synth sound
 1030 matching task with 14 continuous θ_{synth} parameters (more details in Section 4.4). Uncertainties are
 1031 95% CI for 40 training runs using different random seeds and dataset splits. Due to computational
 1032 limitations, the JTFS method is only trained and evaluated for 4 random seeds.

Method	Loudness $L_1 \downarrow$		Spectral Centroid $L_1 \downarrow$		Spectral Flatness $L_1 \downarrow$	
	Micro	Meso	Micro	Meso	Micro	Meso
JTFS	137 ± 10	158 ± 20	859 ± 68	819 ± 96	1200 ± 87	1090 ± 110
SCRAPL (no θ -IS)	389 ± 41	460 ± 60	1020 ± 100	992 ± 110	1800 ± 170	1990 ± 180
SCRAPL	374 ± 39	377 ± 32	1000 ± 120	1080 ± 120	1750 ± 270	1820 ± 220
MSS Lin.	381 ± 12	2510 ± 480	902 ± 27	2350 ± 310	962 ± 43	3620 ± 680
MSS L+L	492 ± 44	1080 ± 91	928 ± 65	1380 ± 76	916 ± 50	1320 ± 150
MSS Rev.	330 ± 21	808 ± 40	1070 ± 49	1540 ± 53	1390 ± 62	2640 ± 96
MSS Rand.	584 ± 75	1030 ± 89	1200 ± 100	1350 ± 99	1690 ± 120	1950 ± 140

1044
 1045
 1046
 1047 Table 9: Drum decay evaluation results for the unsupervised Roland TR-808 DDSP synth sound
 1048 matching task with 14 continuous θ_{synth} parameters (more details in Section 4.4). Uncertainties are
 1049 95% CI for 40 training runs using different random seeds and dataset splits. Due to computational
 1050 limitations, the JTFS method is only trained and evaluated for 4 random seeds.

Method	Loudness $L_1 \downarrow$		Spectral Centroid $L_1 \downarrow$		Spectral Flatness $L_1 \downarrow$	
	Micro	Meso	Micro	Meso	Micro	Meso
JTFS	315 ± 22	355 ± 110	614 ± 51	617 ± 71	527 ± 31	718 ± 190
SCRAPL (no θ -IS)	1810 ± 190	2210 ± 210	1530 ± 170	1860 ± 170	2620 ± 310	3300 ± 370
SCRAPL	1810 ± 160	1740 ± 170	1490 ± 120	1470 ± 140	2540 ± 290	2480 ± 290
MSS Lin.	357 ± 12	1120 ± 260	654 ± 18	1110 ± 160	472 ± 17	1500 ± 350
MSS L+L	389 ± 42	466 ± 45	563 ± 22	597 ± 24	565 ± 29	644 ± 51
MSS Rev.	279 ± 12	494 ± 22	589 ± 21	801 ± 29	552 ± 22	846 ± 29
MSS Rand.	453 ± 21	485 ± 24	660 ± 27	640 ± 35	594 ± 30	658 ± 33

1061
 1062
 1063
 1064
 1065
 1066
 1067
 1068
 1069
 1070
 1071
 1072
 1073
 1074
 1075
 1076
 1077
 1078
 1079

1080 **E EXPERIMENT TRAINING DETAILS AND HYPERPARAMETERS**
 1081
 1082
 1083
 1084

Table 10: Unsupervised granular synth sound matching task hyperparameters.

Category	Hyperparameter Name	Value
Data	N (# of examples)	5120
	train / val / test split	60% / 20% / 20%
Encoder	# of parameters	604 K
	CQT # of octaves	5
	CQT bins / octave	12
	CQT hop length	256
	CQT postprocessing	\log_{10}
	CNN # of conv. blocks	5
	CNN kernel size	$5 \times (3, 3)$
	CNN stride	$5 \times (1, 1)$
	CNN pooling	$5 \times (2, 2)$
	CNN conv. block channels	128
	CNN activation function	PReLU
	CNN embedding dim.	64
	CNN dense layer dropout prob.	0.5
Decoder (Synth)	$\dim(\theta_{\text{synth}})$	2
	sampling rate	8192 Hz
	T (# of samples)	32768
	max. # of grains	64
	grain # of samples	4096
	min. grain pitch	256 Hz
	max. grain pitch	2048 Hz
SCRAPL & JTFS	J	12
	Q_1	8
	Q_2	2
	J_{fr}	3
	Q_{fr}	2
	T	4096
	F	8
	ρ	identity function
	P (# of paths)	315
θ -Importance Sampling	N_{IS} (# of examples)	320
	# of deflated power iterations	20
Training	# of random seed training runs	20
	epochs	200
	batch size	32
	starting learning rate	1×10^{-5}
	learning rate scheduler	none
	Adam β_1	0.9
	Adam β_2	0.999
	weight decay	0.01

1126
 1127
 1128
 1129
 1130
 1131
 1132
 1133

Table 11: Unsupervised AM/FM chirplet synth sound matching task hyperparameters.

Category	Hyperparameter Name	Value
Data	N (# of examples)	5120
	train / val / test split	60% / 20% / 20%
Encoder	# of parameters	604 K
	CQT # of octaves	5
	CQT bins / octave	12
	CQT hop length	256
	CQT postprocessing	\log_{10}
	CNN # of conv. blocks	5
	CNN kernel size	$5 \times (3, 3)$
	CNN stride	$5 \times (1, 1)$
	CNN pooling	$5 \times (2, 2)$
	CNN conv. block channels	128
	CNN activation function	PReLU
	CNN embedding dim.	64
	CNN dense layer dropout prob.	0.5
Decoder (Synth)	$\dim(\theta_{\text{synth}})$	2
	sampling rate	8192 Hz
	T (# of samples)	32768
	chirplet center frequency	512 Hz
	chirplet bandwidth	2 octaves
	min. time shift	-2048 samples
	max. time shift	+2048 samples
SCRAPL & JTFS	J	12
	Q_1	8
	Q_2	2
	J_{fr}	3
	Q_{fr}	2
	T	4096
	F	8
	ρ	identity function
	P (# of paths)	315
θ -Importance Sampling	N_{IS} (# of examples)	32
	# of deflated power iterations	20
Training	# of random seed training runs	20
	epochs	50
	batch size	32
	starting learning rate	1×10^{-4}
	learning rate scheduler	none
	Adam β_1	0.9
	Adam β_2	0.999
	weight decay	0.01

1134

1135

1136

1137

1138

1139

1140

1141

1142

1143

1144

1145

1146

1147

1148

1149

1150

1151

1152

1153

1154

1155

1156

1157

1158

1159

1160

1161

1162

1163

1164

1165

1166

1167

1168

1169

1170

1171

1172

1173

1174

1175

1176

1177

1178

1179

1180

1181

1182

1183

1184

1185

1186

1187

1188
 1189
 1190
 1191
 1192
 1193
 1194
 1195
 1196
 1197
 1198
 1199
 1200
 1201
 1202
 1203
 1204
 1205
 1206
 1207
 1208
 1209
 1210
 1211
 1212
 1213
 1214
 1215
 1216
 1217
 1218
 1219
 1220
 1221
 1222
 1223
 1224
 1225
 1226
 1227
 1228
 1229
 1230
 1231
 1232
 1233
 1234
 1235
 1236
 1237
 1238
 1239
 1240
 1241

Table 12: Unsupervised Roland TR-808 synth sound matching task hyperparameters.

Category	Hyperparameter Name	Value
Data	N (# of examples)	681
	$N_{\text{bass drum}}$	215
	N_{snare}	240
	N_{tom}	189
	$N_{\text{hi-hat}}$	37
	N_{train}	425
	N_{val}	128
	N_{test}	128
Encoder	# of parameters	724 K
	CQT # of octaves	9
	CQT bins / octave	12
	CQT hop length	256
	CQT postprocessing	\log_{10} p
	CNN # of conv. blocks	5
	CNN kernel size	$5 \times (3, 3)$
	CNN stride	$5 \times (1, 1)$
	CNN pooling	(2, 2), (2, 2), (2, 3), (3, 3), (3, 3)
	CNN conv. block channels	128
	CNN activation function	PReLU
	CNN embedding dim.	128
	CNN dense layer dropout prob.	0.25
Decoder (Synth)	$\dim(\theta_{\text{synth}})$	14
	sampling rate	44100 Hz
	T (# of samples)	44100
	min. time shift	-2048 samples
	max. time shift	+2048 samples
SCRAPL & JTFS	J	12
	Q_1	8
	Q_2	2
	J_{fr}	5
	Q_{fr}	2
	T	2048
	F	1
	ρ	\log_{10} p
	P (# of paths)	483
θ -Importance Sampling	N_{IS} (# of examples)	16
	# of deflated power iterations	20
Training	# of random seed training runs	40
	epochs	50
	batch size	8
	starting learning rate	1×10^{-5}
	learning rate scheduler	linearly decreasing until 1×10^{-4}
	Adam β_1	0.9
	Adam β_2	0.999
	weight decay	0.01